



You are here

[Home](#) » [Blogs](#) » [rleishman's blog](#)

Indexing an Oracle Data Warehouse

Submitted by rleishman on Wed, 2007-06-06 03:01

articles:

[Warehousing](#)

Oracle-Bashing

Aside from a nine month excursion to Sybase IQ, I've spent my entire career working with Oracle, so I don't profess too much expertise - indeed any! - about other RDBMS technologies. So in a weak attempt at self-education, I recently accepted an invitation to listen to a Teradata presentation directed at application developers.

Disappointingly, it was more of a sales pitch than a technical decomposition of the benefits of the product. I've long since learned to ignore the promised benefits in these types of pitches, but I love to hear them slag-off the competition; the fastest way to make your sand castle bigger than the other kids' is not to build yours up but to tear theirs down! There are often kernels of truth in these critiques, so I was interested to hear what they had to say about Oracle. They came up with a number of them, but a particular piqued my interest:

- Oracle's parallel query engine is sub-optimal because it is too easy to skew the data across the parallel servers. eg. If roughly clustered in chronological order, then a full table scan for a particular date range - where contiguous chunks are farmed out to each parallel server - will result in some chunks returning no rows of interest and some returning the millions. It's probably easily mitigated with hash-sub-partitioning, but still an interesting concept (for discussion in a perhaps) that is very probably true of many data warehouses.
- Oracle is great for canned reports that use predictable access-paths that can be indexed, but it sucks at high-volume ad analytics and data-mining because you can't index for every possible access path.

What the...??? Did you bother to read the Oracle Data Warehousing manual before you made that one up?

Expekt the Unexpected

This guy's position was that Fact tables in a data warehouse contain a lot of dimensional attributes of which you may want a combination in a query. eg. Widgets sold by Jones in 2007, Widgets sold to ACME in 2007, Jones' sales to ACME in 2007. speaking, a fact table with N dimensions could have $N!$ (ie. $N \times (N-1) \times \dots \times 2 \times 1$) different access paths, requiring $N!$ index this is not tenable beyond $N=4$.

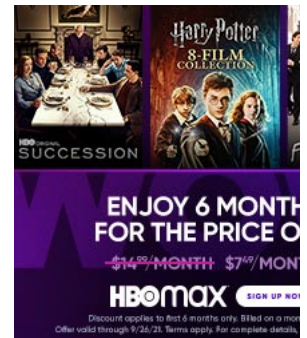
What he'd cheerfully ignored (or at least the research team had cheerfully ignored) is that bitmap indexes solved this issue for queries over 10 years ago in v7.3, partitioning made it scalable in v8.0, and Star Transformations extended the solution to star queries in v8i.

It's hard to get too incensed at the ignorance of a non-Oracle guy (although I admit I gave it a fair shot), because in one way point: although Oracle solved the problem ages ago, a lot of people didn't listen. Oracle data warehouses and data marts are developed today with an archaic indexing strategy that will not scale for ad-hoc analytics.

Layers in the data warehouse

Most data warehouses are built in a layered architecture. Each layer has a different purpose and contains a different representation of data. For this reason, each layer has different characteristics and requires a different indexing strategy.

- Staging Layer** - A place to collect data from source systems before it is transformed. Many ETL implementations use files for the staging layer, others use database tables. Staging tables are truncated before each load and every row is truncated usually in a single full table scan.
- Enterprise Layer** - Also called the 3NF layer, Integrated Data Store, Atomic Data Store, and probably many others. Data from source systems is transformed from the various staging tables into a single 3rd Normal Form enterprise-wide data model. Conventional wisdom is that it is a bad idea to permit user-queries and reporting against the Enterprise Layer; let's not argue just here though. The characteristic usage of the Enterprise Layer is that most access is via a predictable, optimized ETL.
- Presentation Layer** - Also known as the Dimensional Data Store or Star Schema model. Data is denormalised from the Enterprise Layer into the star-schemas of the Presentation Layer, which is then used as the platform to supply data to Business Intelligence (BI) tools. There are two important classes of BI tools: those that access the database interactively and those that access the database via a BI tool.



Site search

User login

Username *
Password *

- [Request new password](#)

Site navigation

- [About](#)
- [Blogs](#)
- [Feed aggregator](#)
- [Books](#)
- [Events](#)
- [FAQ's](#)
- [Forums](#)
- [Mailing Lists](#)
- [Papers](#)
- [Scripts](#)
- [Tools](#)
- [USENET News](#)
- [Wiki](#)
- [XML Feeds](#)



Indexing the Staging Layer

Don't!

The typical life-cycle of data in a Staging table is that it is bulk-loaded, read once in full and transformed into the Enterprise truncated, then the cycle repeats. The load will only be slowed down by the presence of indexes, unique and PK constraints enforced until data gets to the Enterprise Layer, and the transformation will use a Full Table Scan (FTS); there is simply no indexes.

In fact, use Externally Organized Tables (EOTs) for the Staging layer and remove any temptation to index or apply constraints.

Indexing the Enterprise Layer

The Enterprise layer should not be available for user-queries. The purpose of indexing the Enterprise Layer is therefore to:

- Support the transformation of data from the Staging Layer.
- Support the transformation of data to the Presentation Layer.
- Enforce constraints

Since the ETL is the only "user" with access to the Enterprise Layer, there are no unexpected or unknown queries that need handled.

Primary and Unique Keys

Primary and Unique keys should be enforced in the Enterprise Layer, so they must be indexed. For performance reasons, it add additional columns to the index or to make the index non-unique whilst still enforcing the constraint. Take care: this can unexpected side-effects (out of scope for this article). If you want to play it safe, stick to unique indexes enforcing these constraints.

Foreign Keys

Unlike an OLTP system, it is not a hard-and-fast requirement to index every Foreign Key. Indexes are only required to support constraint for deletes on the parent table. Since deletes are typically banned in a data warehouse, an index is not strictly required.

Join Keys

Depending on your ETL tool, transforming data from the Staging Layer may involve joins to the Enterprise Layer. This is a create a non-unique index for every join-key used by the ETL that's not already indexed. It's easy to over-think this: since you scan the entire Staging table, it will often be more efficient for Oracle to FTS the Enterprise tables and perform a hash join - no index required! Let Oracle be the judge; that's what the Cost Based Optimizer (CBO) is for. However if the staging table is small and the Enterprise table is large, the index will be useful, so just create it, gather statistics, and forget about it.

The same argument applies to transformations from Enterprise to Enterprise, Enterprise to Presentation, or Presentation to Enterprise. Index the join-keys on both sides of the join so that the CBO can choose the best join-order.

Filtering Keys

Filtering is uncommon in the Enterprise Layer; most tables are joined by the ETL in-full without filter criteria. In the unlikely event the ETL contains a *selective* WHERE clause (filters out >90%+ of the rows), then index it as for a join-key.

Bitmap and Bitmap-Join Indexes

The Enterprise Layer should not contain any bitmap indexes. They are generally only effective on ad-hoc queries that could benefit from a number of disparate AND and OR conditions on poor-cardinality columns. Where the queries are known and optimised - as in an ETL - a b-tree index will almost always be better.

Indexing the Presentation Layer

For Presentation Layer tables that are only used to build external OLAP cubes, few if any small volume indexed queries will be performed. Index join-keys as for the Enterprise Layer. For databases accessed directly and interactively by the BI tool, follow the below.

Primary and Unique Keys

It is not necessary to enforce constraints in the Presentation Layer since they are enforced in the Enterprise Layer. Tables that are incrementally refreshed from the Enterprise layer will obviously need a Primary Key index to support updates and deletes. A unique key may also be required to support downstream processing, such as replication.

Unique key indexes need not be created.

Primary key indexes should be created for all dimensions to support star-schema joins, but they are optional for Fact tables and are required to support the ETL.

Foreign Keys

Foreign keys in Star-Schema Fact tables should be Bitmap Indexed where the foreign key column contains fewer than 2500 distinct values. Note that this is not the same as a foreign key referencing a Dimension table with 2500 rows, since many rows in the Fact table may not be referenced by the Fact table.

For foreign keys with between 2500 and 10000 distinct values (or those likely to grow beyond 2500) you should trial a Bitmap index. For foreign keys with over 10000 distinct values are unlikely to add benefit - use a B-Tree index.

Recent blog posts

- [Apriori Algorithm in sql.pl/sql and spark sql](#)
- [DB 21c available for download](#)
- [ORDS 21.x make sure you have the latest](#)
- [Time Series Analysis in Spark SQL](#)
- [oracle queries needed](#)
- [Decision Tree Algorithm in Spark SQL](#)
- [database block size - does it really matter?](#)
- [Installing database 19c on Oracle Linux 8](#)
- [Controlling distributed queries with hints](#)
- [Devart Presents Dramatically Improved dbForge Studio for MySQL v.9.0](#)

[More](#)



- Oracle can combine the results of many bitmap index scans. This means you can include any combination of filters on indexed columns and obtain an efficient indexed access using all of those filters.
- In combination with `STAR_TRANSFORMATION`, a query may simply join bitmap indexed foreign keys to dimension table on the dimensions' attributes; Oracle will still use the bitmap indexes in combination as if the filters were applied directly on foreign-key columns.

Join Keys

The only joins in a star-schema should be on foreign keys - see the section above. If joins are required between Facts, then the Presentation layer has been improperly designed - refer to a good dimensional data modelling book.

Snowflake schemas are generally to be avoided, but they do have their uses. A snowflake schema will require primary/foreign key relationships between dimension tables. The primary key of the parent table should already be indexed according to the rules above. The foreign key of the child table should be b-tree indexed to support joins.

Filtering Keys

Filters will be applied mostly to Dimension table columns and occasionally to Fact table columns. A dimension table will seldom be joined in a query to display one or more dimensional attributes, but sometimes it will be used solely as a filter on a single column.

```
SELECT t.txn_dt, sum(t.txn_amt)
FROM   sales_txn t
JOIN   department d USING (dept_key)
WHERE  d.region = 'WEST'
```

To support ad-hoc star-schema queries, it makes sense to index those dimension columns likely to be used as filter keys (eg. `department.region` in the above example). However, to support the type of query shown above, the primary key (department.dept_key) should be appended to the index. In queries such as these, the only columns required from the dimension table are the filter key (for filtering) and the primary key (for joining); if both are available in the index then Oracle does not need to scan the table row at all.

This may seem like overkill; a conformed dimension may contain dozens of columns, any or all of which may be used as filter keys. Well, that's life in a data warehouse. If you want to independently (and efficiently) filter on dozens of different columns then you need dozens of indexes. The additional cost to the ETL is outweighed by the performance benefit to the BI layer.

This approach may be relaxed for small-medium dimensions (say, <5000 rows). Full scanning a table this size will be so fast that the few resources that indexes are not really required.

Filter keys in a fact table should be indexed in the same way as foreign keys.

Indexing Partitioned Tables

When creating an index on a partitioned table, you need to decide whether to partition the index as well. An index that shares its partitioning with its table is termed *locally partitioned*. Alternatively, an index may be non-partitioned (termed: *global non-partitioned*) or partitioned on another column (termed: *globally partitioned*).

Locally partitioned indexes are preferred because they are much easier to manage; partition maintenance activities such as `ALTER TABLE ... DROP ... EXCHANGE PARTITION` can be time and resource consuming on global indexes. Unfortunately, locally partitioned indexes can also be inefficient on queries that do not include a filter on the partition key; the cost of scanning every partition for a small number of rows (say, <10) can be many times slower than a global non-partitioned index (eg. scanning 100 partitions could be up to 100 times slower than a global non-partitioned index for a small number of rows).

The following guidelines will help determine the best index-partitioning strategy for tables in either the Enterprise or Presentation Layer.

- If the primary key does not contain the partition key, use a global index. This is commonplace in atomic (non-aggregated) facts, where the primary key is a transaction identifier and the partition key is a date. Do not blindly add the partition key to the primary key just to avoid a global index: this corrupts the integrity of the primary key.

For very large tables, consider globally partitioning the primary key index on a leading subset of the index columns for easier maintenance.

- Locally partition all bitmap indexes.

Bitmap indexes are typically used individually to identify a large number of rows (and collectively to identify a small number of rows) so they do not suffer the locally partitioned performance issue described above.

- Consider globally indexing alternate keys and filter keys in Enterprise tables that are used to identify a small number of rows.
- Consider globally indexing join keys between two partitioned tables (ie. Where two partitioned tables are joined on a column including the partition key). However if such joins are performed in large volumes (eg. >50,000 rows joined) then you should strongly consider hash-sub-partitioning both tables on the join key in favour of indexing.
- Locally partition all other indexes.

Star Transformation

Lastly, ensure that the initialisation parameter `STAR_TRANSFORMATION_ENABLED` is set to `TRUE` for the Presentation Layer. Without this Oracle will not be able to effectively use the bitmap indexes on star-schema join queries. This is described further in Oracle's Data Warehouse manual.

Hybrid Cases

Not all warehouses are designed in this way. I have seen many cases where a true Enterprise Layer is skipped; the Staging Layer is transformed directly into a dimensional star schema. In such cases, tables must be indexed to optimise both ad-hoc queries and



conflict exists (eg. Bitmap index for a Presentation foreign-key vs. b-tree index for an Enterprise join-key), the Presentation precedence.

Bitmap-Join Indexes

Bitmap-Join indexes provide the ability to index a fact table on a column from a related dimension. This provides a similar efficient result to a regular bitmap index on the foreign key when used in combination with `STAR_TRANSFORMATION`. Consider the case where a dimension has 20 attributes that you want to index; using bitmap-join indexes, you will have 20 indexes on the fact table (a significant overhead on data loads) rather than a single bitmap index on the foreign key.

Like b-tree indexes they can be more efficient than bitmap indexes for any given query, but there does not seem to be an efficient technique to make them adapt to the myriads of different access paths required for ad-hoc queries. Although not specifically Oracle manuals, Bitmap-Join indexes seem to be designed primarily for Oracle OLAP software.

Case Study 1 - Dimension Transformation

The examples in the case study below utilise table structures adapted from Oracle's *Sales History* sample schema.

Table Definitions

Staging	Enterprise	Presentation
SD_ABC_CUSTOMERS	ED_CUSTOMERS	PD_CUSTOMERS
CUST_NUMBER	CUST_ID	CUST_ID
CUST_FIRST_NAME	CUST_NUMBER	CUST_NUMBER
CUST_LAST_NAME	CUST_FIRST_NAME	CUST_FIRST_NAME
CUST_GENDER	CUST_LAST_NAME	CUST_LAST_NAME
CUST_DATE_OF_BIRTH	CUST_GENDER	CUST_GENDER
CUST_MARITAL_STATUS	CUST_YEAR_OF_BIRTH	CUST_YEAR_OF_BIRTH
CUST_STREET_ADDRESS	CUST_MARITAL_STATUS	CUST_MARITAL_STATUS
CUST_POSTAL_CODE	CUST_STREET_ADDRESS	CUST_STREET_ADDRESS
CUST_CITY	CUST_POSTAL_CODE	CUST_POSTAL_CODE
CUST_STATE_PROVINCE	CUST_CITY_ID	CUST_CITY
CUST_COUNTRY	CUST_MAIN_PHONE_NUMBER	CUST_CITY_ID
CUST_MAIN_PHONE_NUMBER	CUST_INCOME_LEVEL	CUST_STATE_PROVINCE
	CUST_CREDIT_LIMIT	CUST_STATE_PROVINCE_ID
	CUST_EMAIL	CUST_COUNTRY_ID
	CUST_SRC_ID	CUST_MAIN_PHONE_NUMBER
	CUST_EFF_FROM	CUST_INCOME_LEVEL
	CUST_EFF_TO	CUST_CREDIT_LIMIT
	CUST_VALID	CUST_EMAIL
		CUST_SRC_ID
		CUST_EFF_FROM
		CUST_EFF_TO
		CUST_VALID

Looking at the structure of these tables, we can see that the Enterprise `ED_CUSTOMERS` is a Type 2 Slowly Changing Dimension. Changes are tracked by inserting a new row and changing the `CUST_EFF_TO` of the previous "current" row. `ED_CUSTOMERS` is populated by many source systems, one of which is **ABC** (the Staging table in this example), which supplies only a subset of columns in the Enterprise table. Different source systems would supply data in a different format and use different staging tables.

The Enterprise layer has been normalised: the city/state/country supplied in the staging table has been transformed into a single `CITY_ID` in the Enterprise table. The Presentation layer shows this structure denormalised in the star-schema dimension `PD_CUSTOMERS`, with city and state attributes added back in.

Lookup Tables

ED_CITIES	ED_STATE_PROVINCES	ED_COUNTRIES
CITY_ID	STATE_PROVINCE_ID	COUNTRY_ID
CITY	STATE_PROVINCE	COUNTRY_ISO_CODE
STATE_PROVINCE_ID	COUNTRY_ID	COUNTRY_NAME
		COUNTRY_SUBREGION_ID
		COUNTRY_NAME_HIST

These tables are used by the Staging to Enterprise ETL in order to transform the textual City/State/Country into a `CITY_ID` foreign key. They are also used in the Enterprise to Presentation ETL to denormalise City and State attributes back into the Customer table.

Sample ETL Code

Staging to Enterprise	Enterprise to Presentation
<pre>SELECT new.* , old.* , cit.city_id FROM sd_abc_customers new LEFT JOIN ed_countries cntr</pre>	<pre>SELECT old.* , new.* , cit.city , stp.state_province , stp.country_id</pre>



```

AND cit.state_province_id = stp.state_province_id ) JOIN    pd_customers new
LEFT JOIN ed_customers old                                ON ( new.cust_id = old.cust_id )
ON ( old.cust_number = new.cust_number
AND old.cust_src_id = 'ABC'
AND old.cust_valid = 'Y' )

```

The above is simply a raw data-retrieval. Additional ETL code would be required to compare new values to old, determine there are any changes, and perform appropriate inserts and updates on the target table.

Required Indexes

Table	Index Columns	Reason / Notes
SD_CUSTOMERS	None!	Staging tables are never indexed
ED_CUSTOMERS	CUST_ID	Enforce constraint - Primary Key
ED_CUSTOMERS	CUST_SRC_ID	Enforce constraint - Unique Key
	CUST_NUMBER	
	CUST_EFF_FROM	
ED_CUSTOMERS	CUST_SRC_ID	Join Key - Enterprise transformation
	CUST_NUMBER	
	CUST_VALID	
ED_COUNTRIES		Assume PK and Unique indexes as for ED_CUSTOMERS.
ED_STATE_PROVINCES		
ED_CITIES		
ED_CUSTOMERS	CUST_CITY_ID	Join Key - Presentation transformation
PD_CUSTOMERS	CUST_ID	Join Key - Presentation transformation (nonunique: not required for prim enforcement)
PD_CUSTOMERS	CUST_COUNTRY_ID	Foreign Key - snowflaked dimension
PD_CUSTOMERS	CUST_NUMBER	Filter Key - Business Intelligence
	CUST_ID	
PD_CUSTOMERS	CUST_LAST_NAME	Filter Key - Business Intelligence
	CUST_ID	
PD_CUSTOMERS	CUST_POSTAL_CODE	Filter Key - Business Intelligence
	CUST_ID	
PD_CUSTOMERS	CUST_CITY	Filter Key - Business Intelligence
	CUST_ID	
PD_CUSTOMERS	CUST_MAIN_PHONE_NUMBER	Filter Key - Business Intelligence
	CUST_ID	

Case Study 2 - Fact Transformation

The examples in the case study below utilise table structures adapted from Oracle's *Sales History* sample schema.

Table Definitions

Staging	Enterprise	Presentation
SF_ABC_SALES_TRANSACTIONS	EF_SALES_TRANSACTIONS	PF_SALES_DAY
TRANSACTION_NUMBER	TRANSACTION_SRC_ID	PROD_ID
PROD_CODE	TRANSACTION_NUMBER	CUST_ID
CUST_NUMBER	PROD_ID	TIME_ID
TRANSACTION_DATE	CUST_ID	CHANNEL_ID
CHANNEL_CODE	TIME_ID	PROMO_ID
PROMO_NUMBER	CHANNEL_ID	QUANTITY_SOLD
QUANTITY_SOLD	PROMO_ID	AMOUNT_SOLD
AMOUNT_SOLD	QUANTITY_SOLD	
	AMOUNT_SOLD	

Lookup Tables

ED_PRODUCTS	ED_CUSTOMERS	ED_TIMES	ED_CHANNELS	ED_PROMOTIONS
PROD_ID	CUST_ID	TIME_ID	CHANNEL_ID	PROMO_ID
PROD_CODE	CUST_NUMBER		CHANNEL_CODE	PROMO_NUMBER
PROD_NAME	CUST_FIRST_NAME		CHANNEL_DESC	PROMO_NAME
PROD_DESC	CUST_LAST_NAME		CHANNEL_CLASS_ID	PROMO_SUBCATEGORY_ID
PROD_SUBCATEGORY_ID	CUST_GENDER			PROMO_COST
PROD_WEIGHT_CLASS	CUST_YEAR_OF_BIRTH			PROMO_BEGIN_DATE
PROD_UNIT_OF_MEASURE	CUST_MARITAL_STATUS			PROMO_END_DATE
PROD_PACK_SIZE	CUST_STREET_ADDRESS			
SUPPLIER_ID	CUST_POSTAL_CODE			
PROD_STATUS	CUST_CITY_ID			
PROD_LIST_PRICE	CUST_MAIN_PHONE_NUMBER			
PROD_MIN_PRICE	CUST_INCOME_LEVEL			
PROD_SRC_ID	CUST_CREDIT_LIMIT			
PROD_EFF_FROM	CUST_EMAIL			



These tables are used by the Staging to Enterprise ETL in order to transform the natural keys (eg. CUSTOMER_NUMBER) into surrogate keys (eg. CUST_ID).

Sample ETL Code

Staging to Enterprise	Enterprise to Presentation
<pre> SELECT old.* , new.* , prod.prod_id , cust.cust_id , tim.time_id , chnl.channel_id , promo.promo_id FROM sf_abc_sales_transactions new LEFT JOIN ed_products prod ON (prod.prod_code = new.prod_code AND prod.prod_src_id = 'ABC' AND new.transaction_date BETWEEN prod.prod_eff_from AND prod.prod_eff_to) LEFT JOIN ed_customers cust ON (cust.cust_number = new.cust_number AND cust.cust_src_id = 'ABC' AND new.transaction_date BETWEEN cust.cust_eff_from AND cust.cust_eff_to) LEFT JOIN ed_times tim ON (tim.time_id = new.transactions_date) LEFT JOIN ed_channels chnl ON (chnl.channel_code = new.channel_code) LEFT JOIN ed_promotions promo ON (promo.promo_number = new.promo_number) LEFT JOIN ef_sales_transactions old ON (old.transaction_number = new.transaction_number AND old.transaction_src_id = 'ABC') </pre>	<pre> SELECT new.* , old.* FROM (SELECT prod_id , cust_id , time_id , channel_id , promo_id , SUM(quantity_sold) AS quantity_sold , SUM(amount_sold) AS amount_sold FROM ef_sales_transactions GROUP BY prod_id , cust_id , time_id , channel_id , promo_id) new LEFT JOIN pf_sales_day old ON (old.prod_id = new.prod_id AND old.cust_id = new.cust_id AND old.time_id = new.time_id AND old.channel_id = new.channel_id AND old.promo_id = new.promo_id) </pre>

The above is simply a raw data-retrieval. Additional ETL code would be required to compare new values to old, determine there are any changes, and perform appropriate inserts and updates on the target table.

Required Indexes

Table	Index Columns	Reason / Notes
SF_ABC_SALES_TRANSACTIONS	None!	Staging tables are never indexed
EF_SALES_TRANSACTIONS	TRANSACTION_SRC_ID TRANSACTION_NUMBER	Enforce constraint - Primary Key
ED_PRODUCTS		Assume PK and join keys indexed as for ED_CUSTOMERS above.
ED_CUSTOMERS		
ED_TIMES		
ED_CHANNELS		
ED_PROMOTIONS		
EF_SALES_TRANSACTIONS	PROD_ID CUST_ID TIME_ID CHANNEL_ID PROMO_ID	Join Key - Presentation Transformation
PF_SALES_DAY	PROD_ID CUST_ID TIME_ID CHANNEL_ID PROMO_ID	Join Key - Presentation Transformation (also PK - not enforced)
PF_SALES_DAY	PROD_ID	Foreign Key - Bitmap Index
PF_SALES_DAY	CUST_ID	Foreign Key - B-Tree Index (>2500 distinct values)
PF_SALES_DAY	TIME_ID	Foreign Key - Bitmap Index
PF_SALES_DAY	CHANNEL_ID	Foreign Key - Bitmap Index
PF_SALES_DAY	PROMO_ID	Foreign Key - Bitmap Index

Conclusion

If you index based on a set of rules (like those above) then you will almost certainly get a sub-optimal result; some indexes used, some will give no improvement, and some useful indexes or variants will be missed. What this method *will* deliver is baseline upon which you may start tuning; all processes and queries should run acceptably fast, and most changes will yield improvements.

»

- [reishman's blog](#)
- [Log in](#) to post comments

:: [Blogger Home](#) :: [Wiki Home](#) :: [Forum Home](#) :: [Privacy](#) :: [Contact](#) ::



