

Tutorial : Getting Started with Kubernetes with Docker on Mac



Romin Irani

Follow

Jan 10, 2018 · 9 min read

If you are looking for running Kubernetes on your Windows laptop, go to [this tutorial](#).

This blog post is related to [Getting Started with Kubernetes on your Windows laptop with Minikube](#) but this time with a Mac machine. The other big difference here is that this is not with Minikube, which you can still install on a Mac. It is with a Edge version of Docker on Mac.



This tutorial works on the Edge version of Docker on Mac and could undergo changes as it approaches a stable release. I will keep the article updated.

We shall cover the following in this post:

- Installing Docker on Mac Edge version
- Go through the basic Kubernetes commands to validate our environment.

This tutorial assumes that you know about Docker and Kubernetes in general. To quote from my previous article, I do not want to spend time explaining about what Kubernetes is and its building blocks like Pods, Replication Controllers, Services, Deployments and more. There are multiple articles on that and I suggest that you go through it.

I have written a couple of other articles that go through a high level overview of Kubernetes:

- [Introduction to Kubernetes](#)
- [Kubernetes Building Blocks](#)

It is important that you go through some basic material on its concepts, so that we can directly get down into its commands.

Docker for Mac installation

As per the [official documentation](#), **Kubernetes is only available in Docker for Mac 17.12 CE Edge**. Go to the [official download page](#) and click on the Edge channel and not the Stable version.

Edge channel

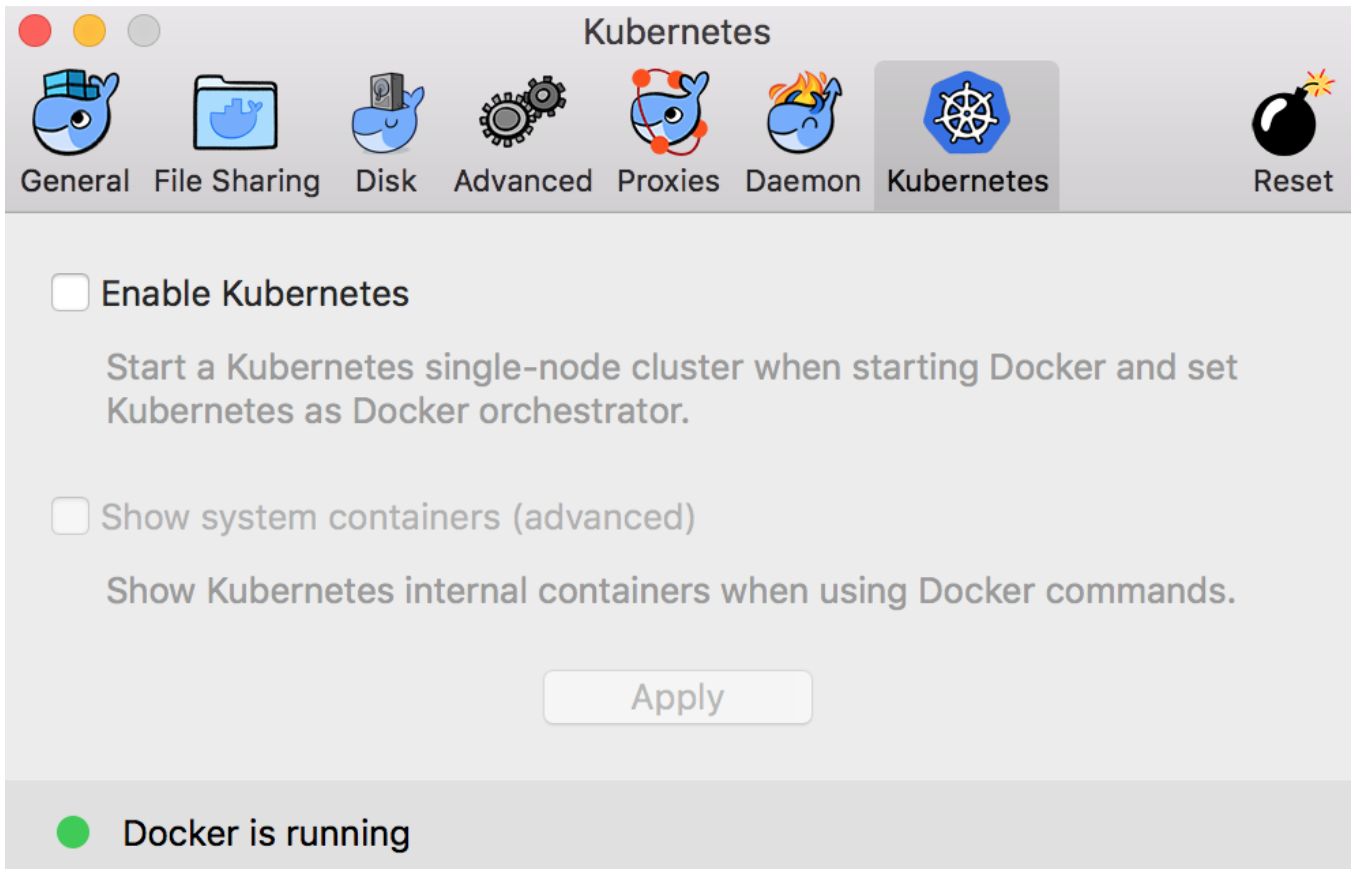
This installer provides the latest Edge release of Docker for Mac and Engine, and typically offers new features in development. Use this channel if you want to get experimental features faster, and can weather some instability and bugs. We collect all usage data on Edge releases across the board.

Edge builds are released once per month.

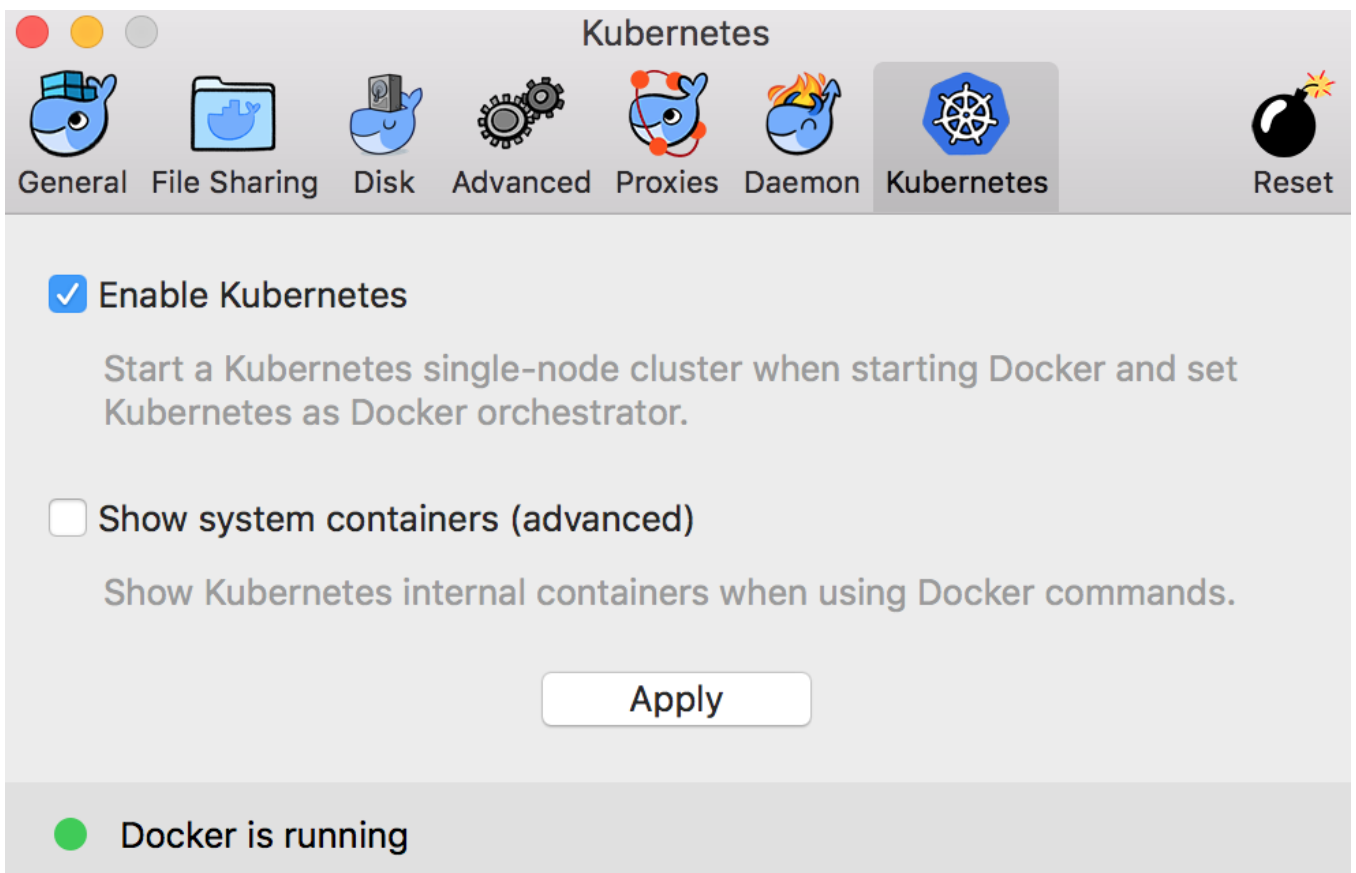
[Get Docker for Mac \(Edge\)](#)

Download the .dmg file and go ahead with the standard installation steps. You can then launch Docker Edge. Click on the Docker icon and go to **Preferences** window as shown

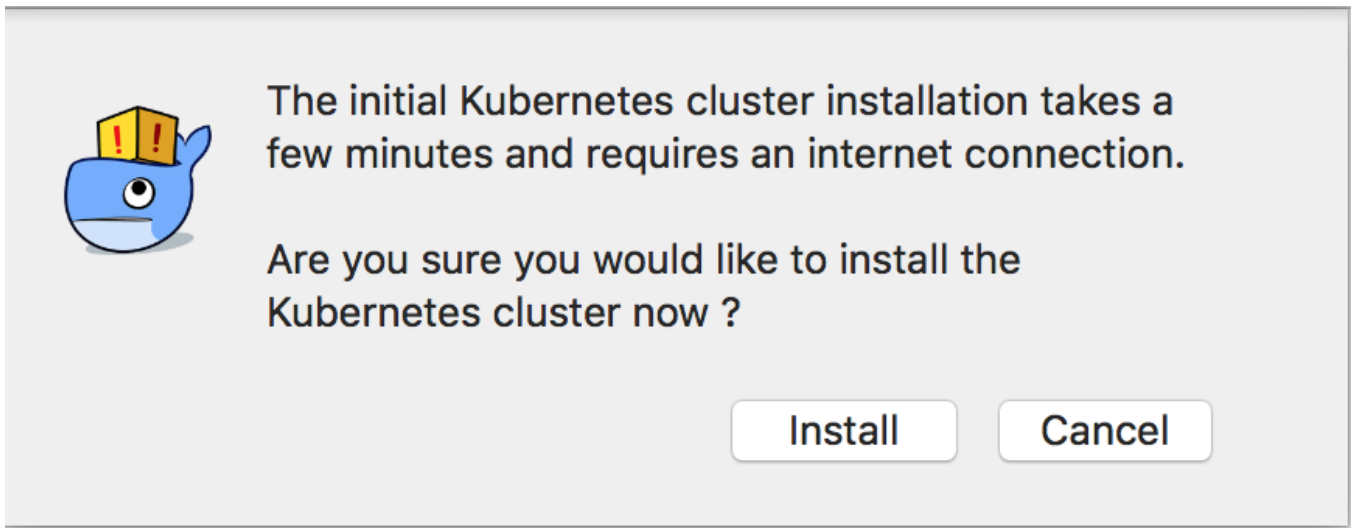
below. Click on the **Kubernetes** icon.



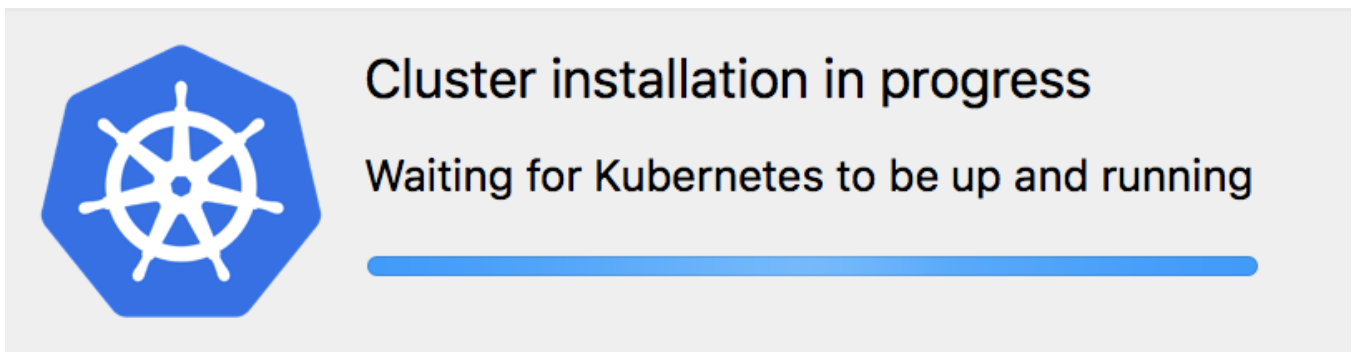
You will notice that Kubernetes is not enabled. Simply check on the **Enable Kubernetes** option and then hit the **Apply** button as shown below:



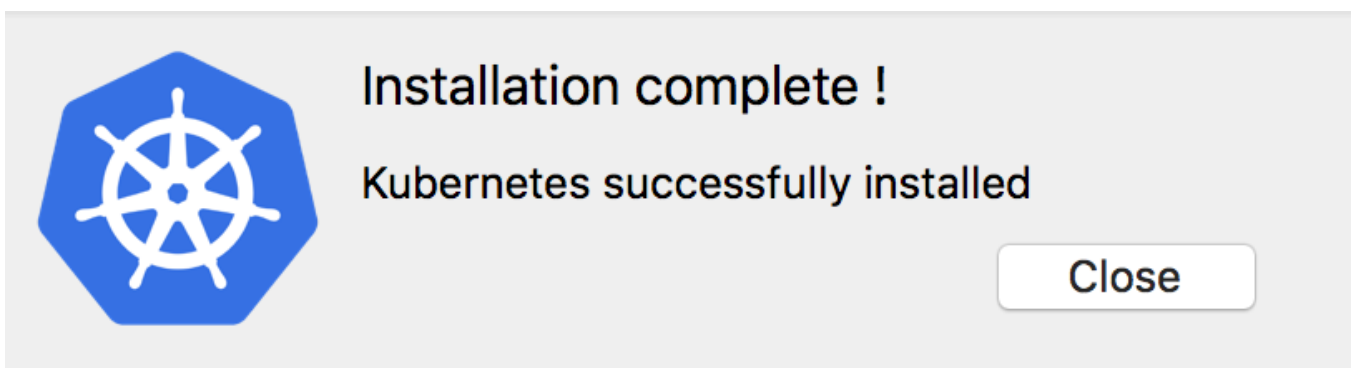
This will display a message that the Kubernetes cluster needs to be installed. Make sure you are connected to the Internet and click on **Install**



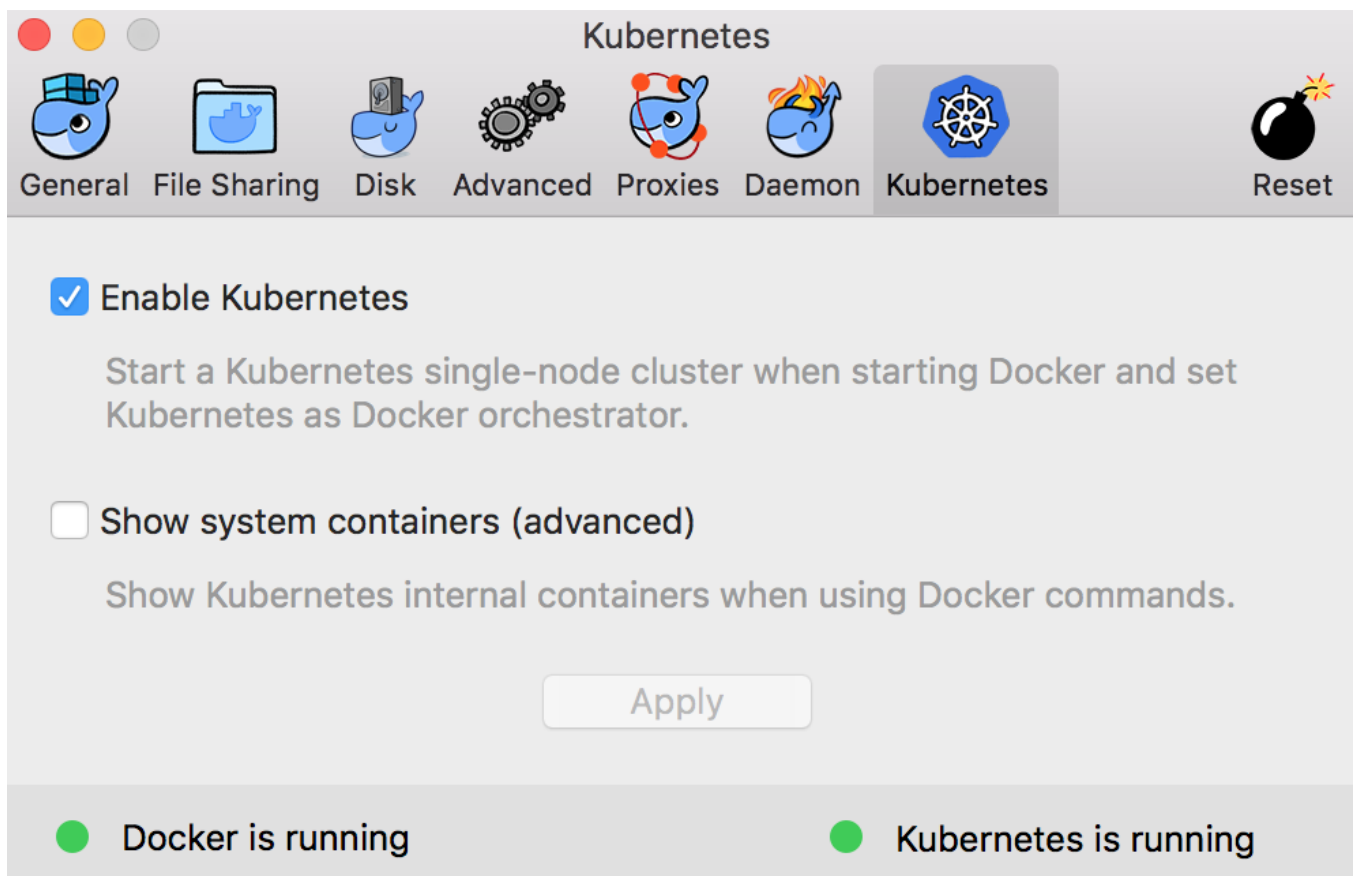
The installation starts. Please be patient since this could take a while depending on your network. *It would have been nice to see a small log window that shows a sequence of steps.*



Finally, you should see the following message:




Click on Close. This will lead you back to the Preferences dialog and you should see the following screen:



Note the two messages at the bottom of the window mentioning:

- Docker is running
- Kubernetes is running

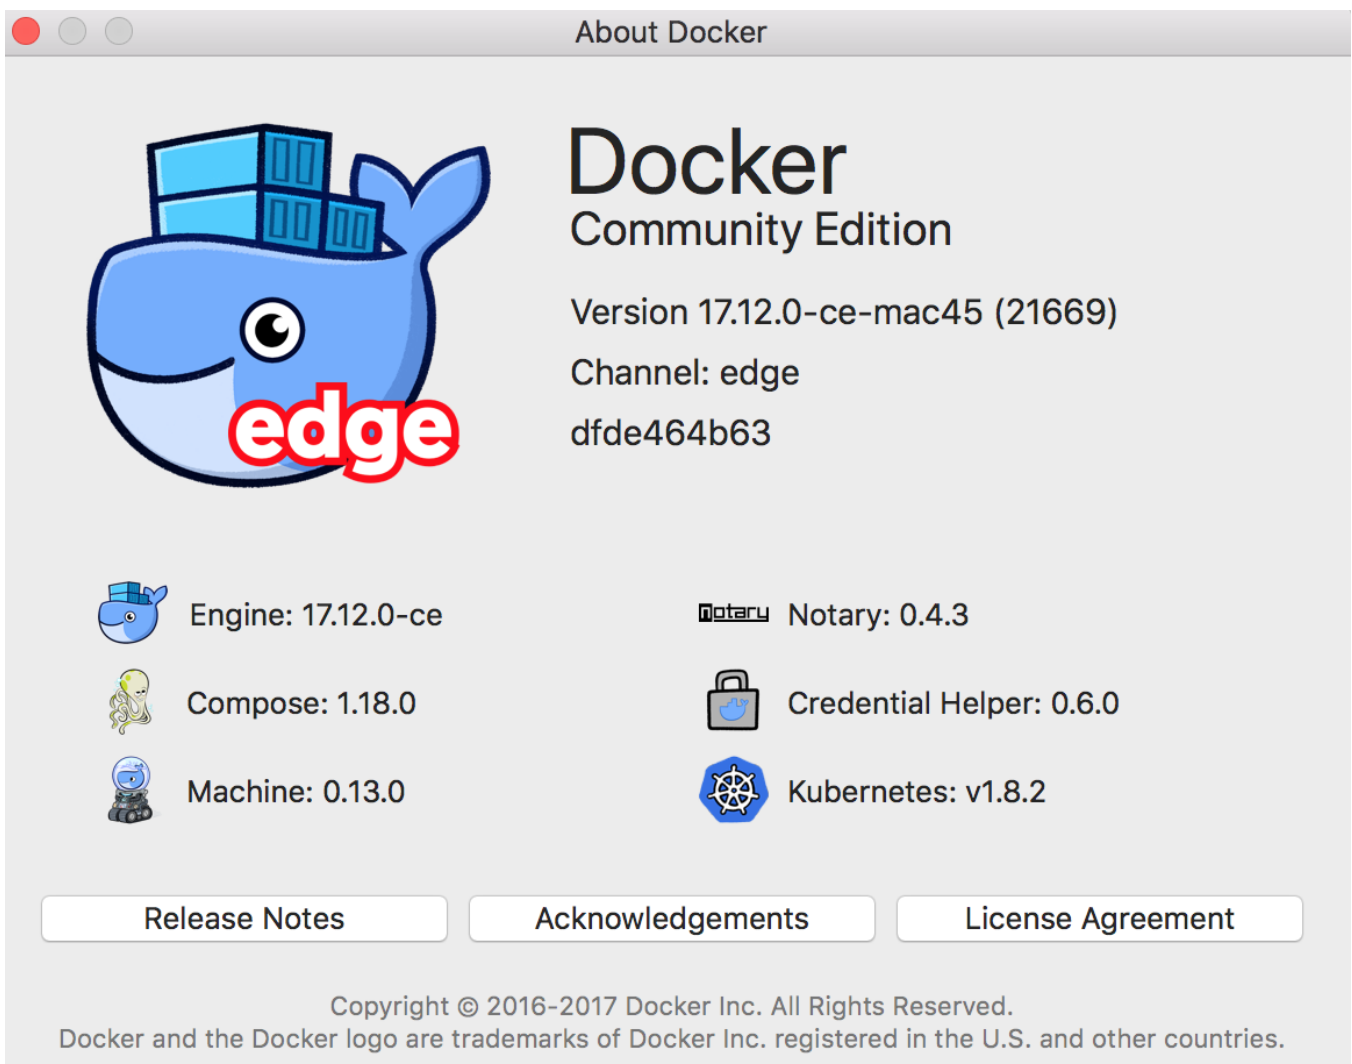
In case you stop running and try to run Docker again, you will notice that both Docker and Kubernetes services are starting as shown below:



Congratulations! You now have the following:

- A standalone Kubernetes server and client, as well as Docker CLI integration.
- The Kubernetes server is a single-node cluster and is not configurable.

Just FYI ... my About Docker shows the following:



Check our installation

Let us try out a few things to ensure that we can make sense of what has got installed. Execute the following commands in a terminal:

```
$ kubectl version
```

```
Client Version: version.Info{Major:"1", Minor:"8",  
GitVersion:"v1.8.4",  
GitCommit:"9befc2b8928a9426501d3bf62f72849d5cbcd5a3",  
GitTreeState:"clean", BuildDate:"2017-11-20T05:28:34Z",  
GoVersion:"go1.8.3", Compiler:"gc", Platform:"darwin/amd64"}
```

```
Server Version: version.Info{Major:"1", Minor:"8",  
GitVersion:"v1.8.2",  
GitCommit:"bdaeafa71f6c7c04636251031f93464384d54963",  
GitTreeState:"clean", BuildDate:"2017-10-24T19:38:10Z",  
GoVersion:"go1.8.3", Compiler:"gc", Platform:"linux/amd64"}
```

You might have noticed that my server and client versions are different. I am using kubectl from my gCloud SDK tools and Docker for Mac, when it launched the

Kubernetes cluster has been able to set the cluster context for the kubectl utility for you. So if we fire the following command:

```
$ kubectl config current-context  
docker-for-desktop
```

You can see that the cluster is set to **docker-for-desktop**.

Tip: In case you switch between different clusters, you can always get back using the following:

```
$ kubectl config use-context docker-for-desktop  
Switched to context "docker-for-desktop"
```

Let us get some information on the cluster.

```
$ kubectl cluster-info
```

```
Kubernetes master is running at https://localhost:6443  
KubeDNS is running at https://localhost:6443/api/v1/namespaces/kube-system/services/kube-dns/proxy
```

Let us check out the nodes in the cluster:

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
docker-for-desktop	Ready	master	7h	v1.8.2

Installing the Kubernetes Dashboard

The next step that we need to do here is to install the Kubernetes Dashboard. We can use the Kubernetes Dashboard YAML that is available and submit the same to the Kubernetes Master as follows:

```
$ kubectl create -f  
https://raw.githubusercontent.com/kubernetes/dashboard/master/src/deploy/recommended/kubernetes-dashboard.yaml
```

```
secret "kubernetes-dashboard-certs" created
serviceaccount "kubernetes-dashboard" created
role "kubernetes-dashboard-minimal" created
rolebinding "kubernetes-dashboard-minimal" created
deployment "kubernetes-dashboard" created
service "kubernetes-dashboard" created
```

The Dashboard application will get deployed as a Pod in the **kube-system** namespace. We can get a list of all our Pods in that namespace via the following command:

```
$ kubectl get pods - namespace=kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
etcd-docker-for-desktop	1/1	Running	0	8h
kube-apiserver-docker-for-desktop	1/1	Running	0	7h
kube-controller-manager-docker-for-desktop	1/1	Running	0	8h
kube-dns-545bc4bfd4-l9tw9	3/3	Running	0	8h
kube-proxy-w8pq7	1/1	Running	0	8h
kube-scheduler-docker-for-desktop	1/1	Running	0	7h
kubernetes-dashboard-7798c48646-ctrl	1/1	Running	0	3m


Ensure that the Pod shown in bold is in Running state. It could take some time to change from **ContainerCreating** to **Running**, so be patient.

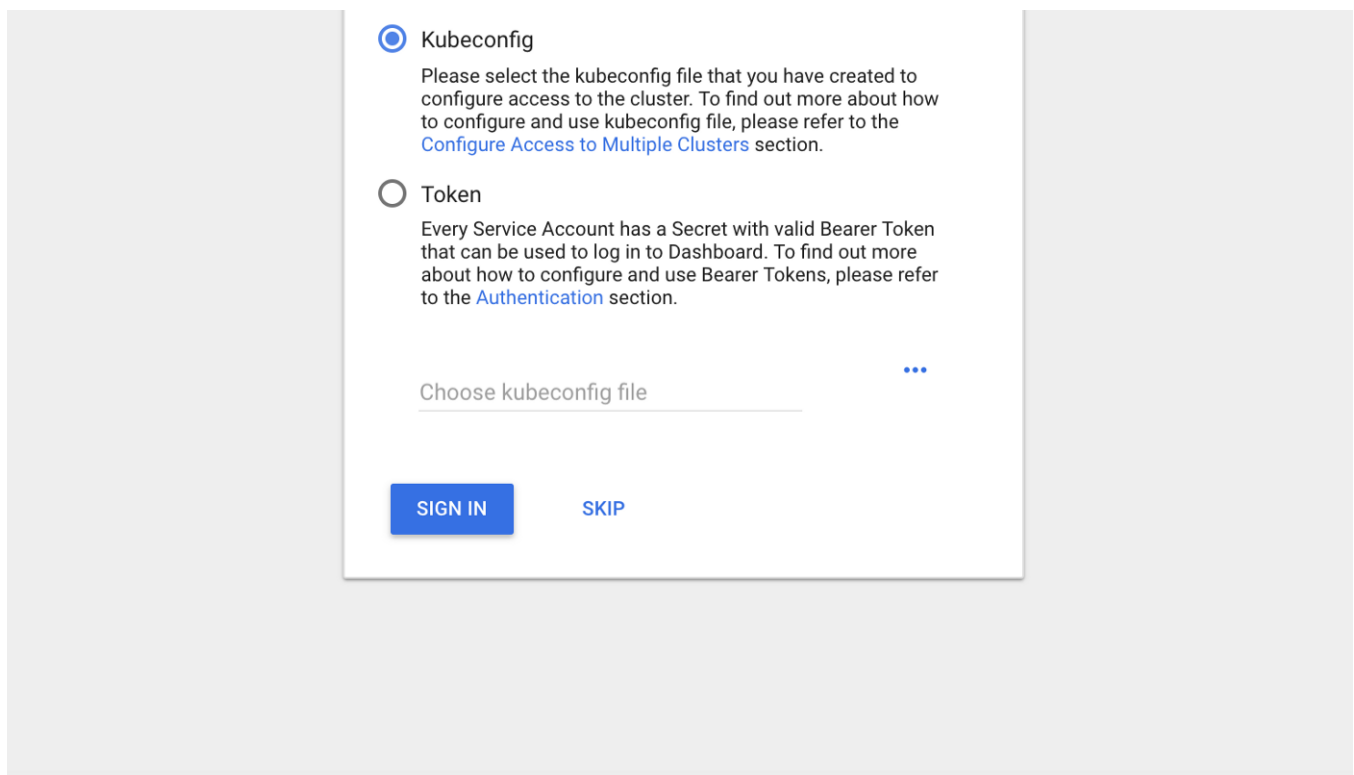
Once it is in running state, you can setup a forwarding port to that specific Pod. So in our case, we can setup 8443 for the Pod Name as shown below:

```
$ kubectl port-forward kubernetes-dashboard-7798c48646-ctrl
8443:8443 - namespace=kube-system
```

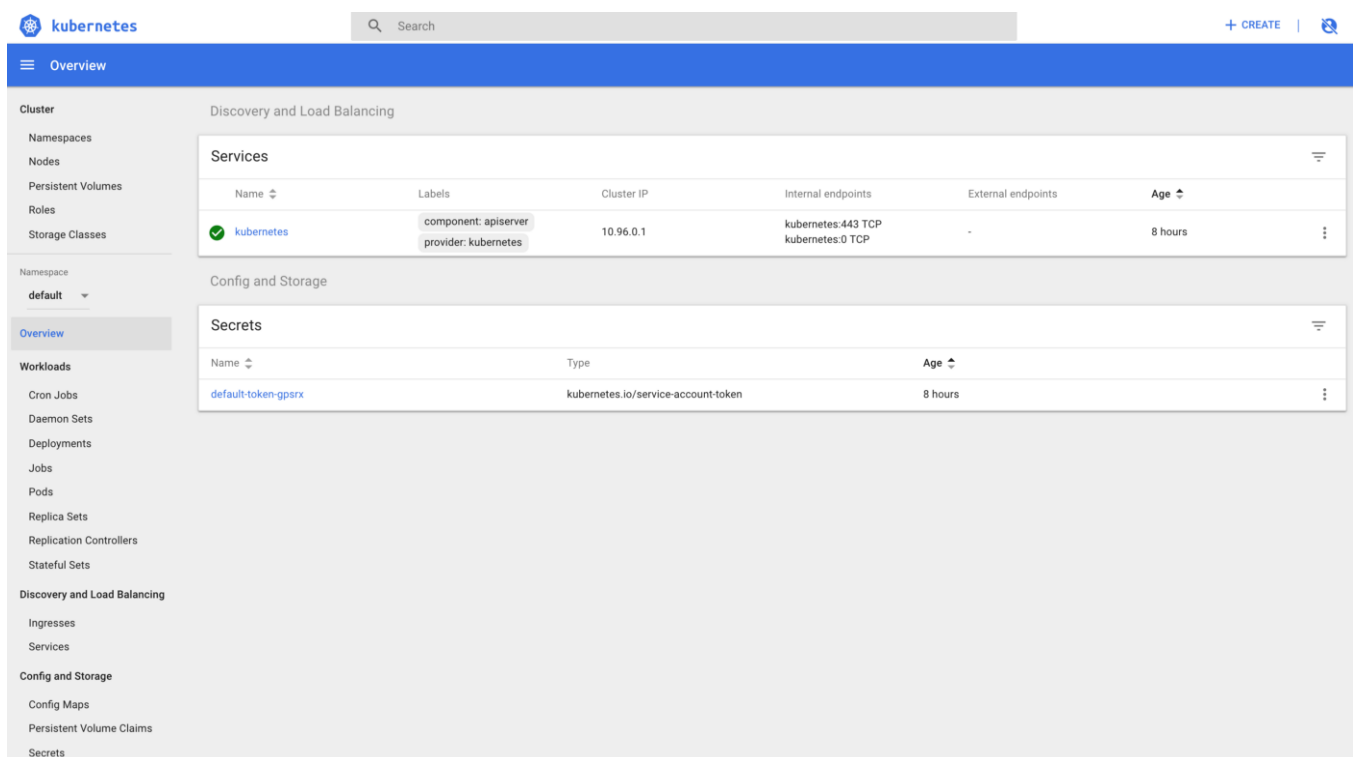
Forwarding from 127.0.0.1:8443 -> 8443

You can now launch a browser and go to <https://localhost:8443>. You might see some warnings but proceed. You will see the following screen:







Click on **SKIP** and you will be lead to the Dashboard as shown below:



Click on Nodes and you will see the single node as given below:



Cluster	
Namespaces	
Nodes	
Persistent Volumes	
Roles	
Storage Classes	
Namespace	

Nodes		
Name	Labels	Ready
 docker-for-desktop	beta.kubernetes.io/arch=amd64 beta.kubernetes.io/os=linux kubernetes.io/hostname=... node-role.kubernetes.io/master=...	True

Running a Workload

Let us proceed now to running a simple Nginx container to see the whole thing in action:

We are going to use the run command as shown below:

```
$ kubectl run hello-nginx --image=nginx --port=80
```

deployment "hello-nginx" created


This creates a deployment and we can investigate into the Pod that gets created, which will run the container:

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-nginx-5d47cdc4b7-wxf9b	0/1	ContainerCreating	0	16s

You can see that the STATUS column value is **ContainerCreating**.

Now, let us go back to the Dashboard and see the Deployments:

Workloads > Deployments				
<div>Workloads</div> <div> Cron Jobs Daemon Sets Deployments Jobs Pods </div>				
Deployments				
Name	Labels	Pods	Age	Images
 hello-nginx	run: hello-nginx	0 / 1	-	nginx

You can notice that if we go to the Deployments option, the Deployment is listed and the status is still in progress. You can also notice that the Pods value is 0/1.

If we wait for a while, the Pod will eventually get created and it will ready as the command below shows:

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-nginx-5d47cdc4b7-wxf9b	1/1	Running	0	3m

Name	Labels	Pods	Age	Images
hello-nginx	run: hello-nginx	1 / 1	a minute	nginx

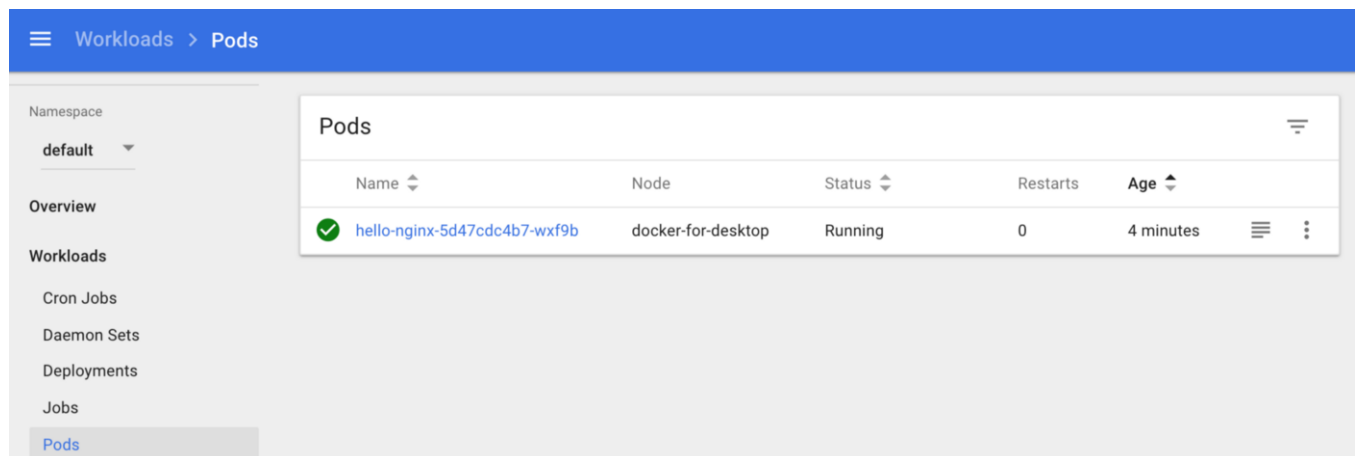
If we visit the Replica Sets now, we can see it:

Name	Labels	Pods	Age	Images
hello-nginx-5d47	pod-template-... run: hello-nginx	1 / 1	3 minutes	nginx

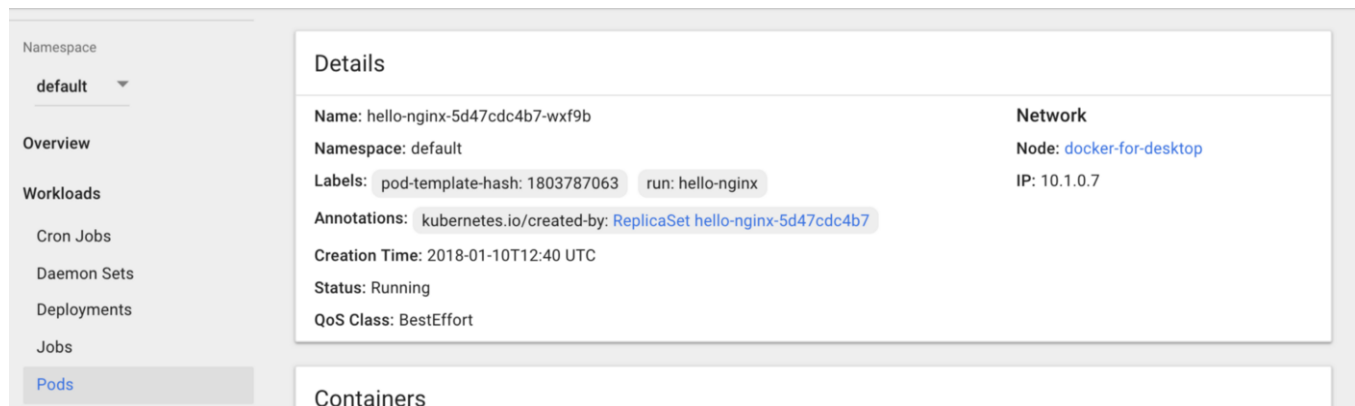
Click on the Replica Set name and it will show the Pod details as given below:

Name	Node	Status	Restarts	Age
hello-nginx-5d47cdc	docker-for-desktop	Running	0	3 minutes

Alternately, you can also get to the Pods via the **Pods** link in the Workloads as shown below:



Click on the Pod and you can get various details on it as given below:



You can see that it has been given some default labels. You can see its IP address. It is part of the node named **docker-for-desktop**.

There are some interesting links that you will find on this page as shown below, via which you can directly EXEC into the pods or see the logs too.



We could have got the Node and Pod details via a variety of **kubectl describe node/pod** commands and we can still do that. An example of that is shown below:

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-nginx-5d47cdc4b7-wxf9b	1/1	Running	0	10m

```
$ kubectl describe pod hello-nginx-5d47cdc4b7-wxf9b
```

```
Name: hello-nginx-5d47cdc4b7-wxf9b
Namespace: default
Node: docker-for-desktop/192.168.65.3
Start Time: Wed, 10 Jan 2018 18:10:35 +0530
Labels: pod-template-hash=1803787063
run=hello-nginx
```

```
Annotations: kubernetes.io/created-by=
{"kind":"SerializedReference","apiVersion":"v1","reference":
{"kind":"ReplicaSet","namespace":"default","name":"hello-nginx-
5d47cdc4b7","uid":"7415cff7-f603-11e7-9f7b-025000000000...
```

```
Status: Running
IP: 10.1.0.7
Created By: ReplicaSet/hello-nginx-5d47cdc4b7
Controlled By: ReplicaSet/hello-nginx-5d47cdc4b7
```

```
Containers:
hello-nginx:
Container ID:
docker://a0c3309b61be4473bf6924ea2be9795de660f49bda36492785f94627690
cbdae
Image: nginx
Image ID: docker-
pullable://nginx@sha256:285b49d42c703fdf257d1e2422765c4ba9d3e37768d6
ea83d7fe2043dad6e63d
Port: 80/TCP
State: Running
```

```
...// REST OF THE OUTPUT
```

Expose a Service

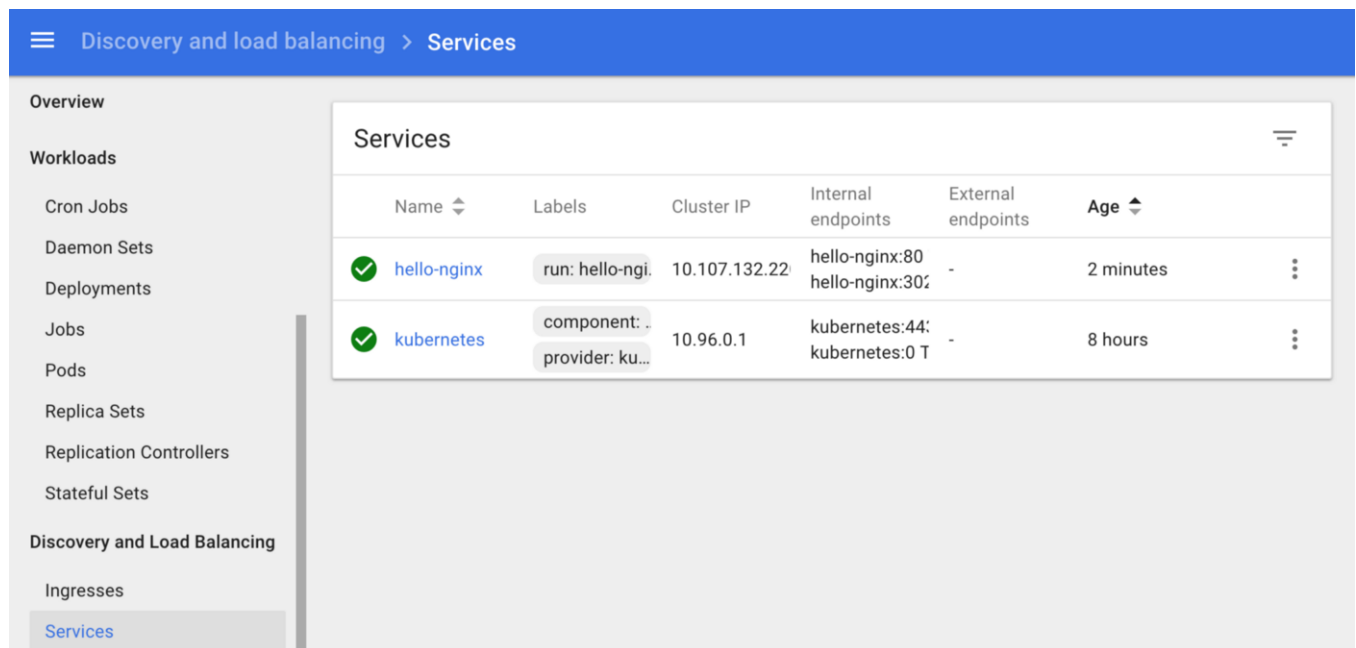
It is time now to expose our basic Nginx deployment as a service. We can use the command shown below:

```
$ kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
hello-nginx	1	1	1	1	19m

```
$ kubectl expose deployment hello-nginx --type=NodePort
service "hello-nginx" exposed
```

If we visit the Dashboard at this point and go to the Services section, we can see out **hello-nginx** service entry.



Alternately, we can use kubectl too, to check it out:

```
$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-nginx	NodePort	10.107.132.220	<none>	80:30259/TCP	1m
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8h

and

```
$ kubectl describe service hello-nginx
```

```
Name: hello-nginx
Namespace: default
Labels: run=hello-nginx
Annotations: <none>
Selector: run=hello-nginx
Type: NodePort
IP: 10.107.132.220
Port: <unset> 80/TCP
TargetPort: 80/TCP
NodePort: <unset> 30259/TCP
Endpoints: 10.1.0.7:80
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
```

Scaling the Service

OK, I am not yet done!

When we created the deployment, we did not mention about the number of instances for our service. So we just had one Pod that was provisioned on the single node.

Let us go and see how we can scale this via the scale command. We want to scale it to 3 Pods.

```
$ kubectl scale --replicas=3 deployment/hello-nginx
deployment "hello-nginx" scaled
```

We can see the status of the deployment in a while:

```
$ kubectl get deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
hello-nginx	3	3	3	3	45m

Now, if we visit the Dashboard for our Deployment:

The screenshot shows the Kubernetes Dashboard with the 'Workloads > Deployments' view selected. A table titled 'Deployments' contains one entry: 'hello-nginx'. The table columns are Name, Labels, Pods, Age, and Images. The 'Pods' column shows '3 / 3' and the 'Age' column shows '44 minutes'. A green checkmark is next to the deployment name.

We have the 3/3 Pods available. Similarly, we can see our Service or Pods.

The screenshot shows the Kubernetes Dashboard with the 'Discovery and load balancing > Services > hello-nginx' view selected. A table titled 'Pods' contains two entries. The table columns are Name, Node, Status, Restarts, and Age. Both pods are in 'Running' status with 0 restarts and an age of 37 seconds. A green checkmark is next to the first pod name.

Pods	✓ hello-nginx-5d47c	desktop	running	0	37 seconds		
Replica Sets	✓ hello-nginx-5d47c	docker-for-desktop	Running	0	45 minutes		
Replication Controllers							
Stateful Sets							
Discovery and Load Balancing							
Ingresses							
Services							

Events						
Message	Source	Sub-object	Count	First seen	Last seen	
Scaled up replica set hello-nginx-5d47c	deployment-nginx	-	1	2018-01-10T10:40:10Z	2018-01-10T10:40:10Z	

Conclusion

Hope this blog post gets you started with Kubernetes with Docker for Mac. Please let me know about your experience in the comments. Now go forth and play the role of a helmsman.

[Kubernetes](#) [Docker For Mac](#) [Docker](#) [Tutorial](#) [K8s](#)

[About](#) [Write](#) [Help](#) [Legal](#)

Get the Medium app

